

ECE 382N-Sec (FA25):

L12: Information-Flow Tracking in Hardware

Neil Zhao

neil.zhao@utexas.edu

Recap: Speculative Taint Tracking (STT)

Tainted 1 void vulnerable_code(size_t idx) {
Root of taint 2 if (idx < array_size) { // B1
3 u8 data = array[idx]; // speculative access, execute & taint!
4 u8 X = data + 2; // not a transmitter, execute!
5 if (idx % 2) { // B2
6 u8 junk = glob_array[X*4096]; // transmitter, delay!
7 }
8 }
9 }



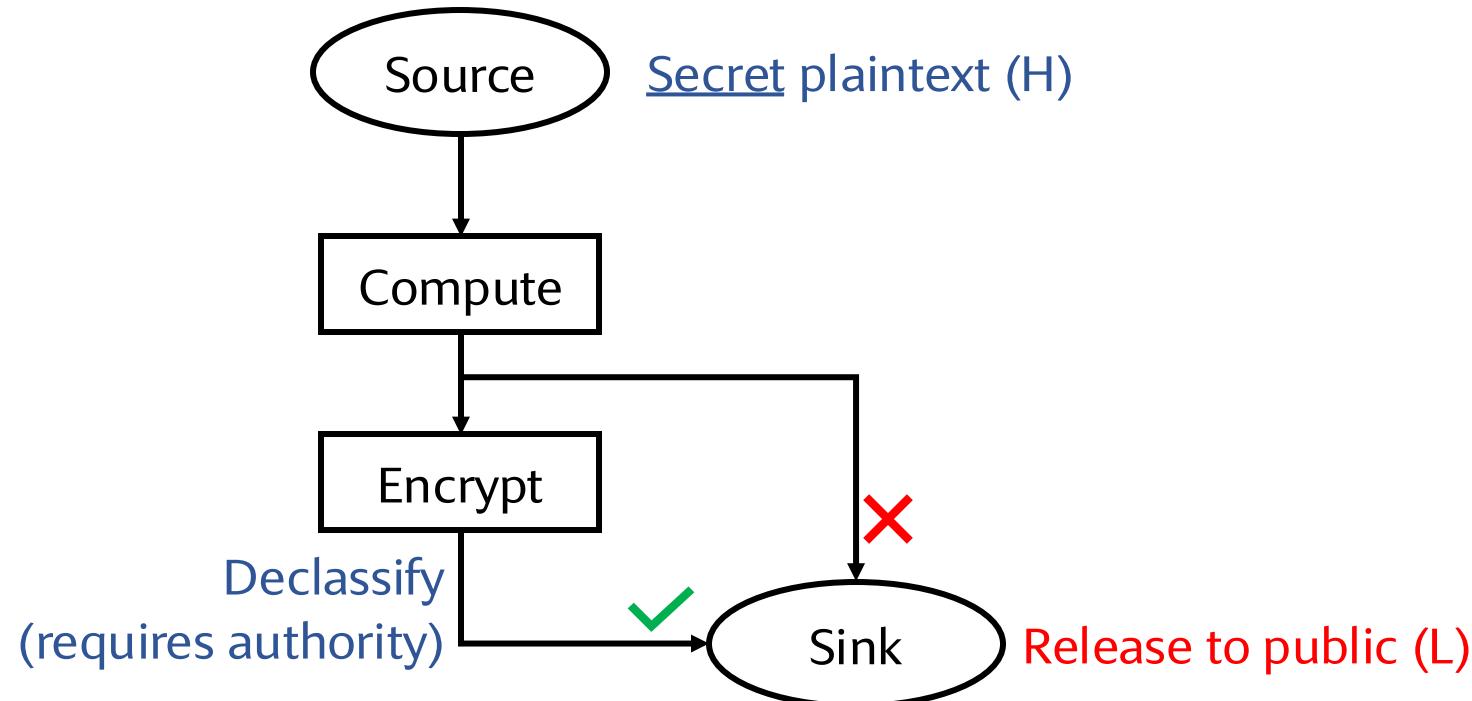
Recap: Speculative Taint Tracking (STT)

Leakage through implicit channels

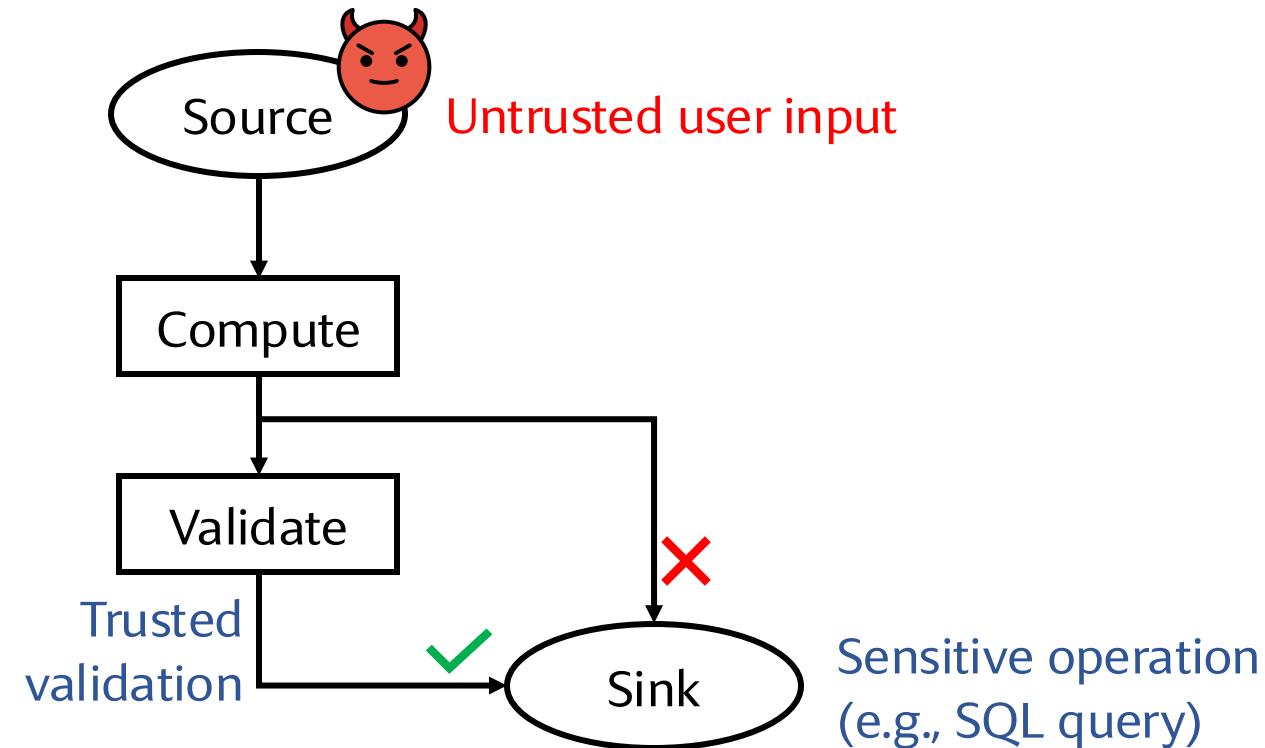
Tainted

```
1 void vulnerable_code(size_t idx) {
2     if (idx < array_size) {
3         u8 data = array[idx]; // speculative access
4         if (data & 1) { // B1
5             div();
6         } else {
7             mul();
8         }
9     }
10 }
```

Information Flow for Confidentiality



Information Flow for Integrity



Explicit Information Flow

```
(H) R1 = load ptr_to_secret;  
(H) R2 = R1 + 10;  
(H) R3 = R2 & R1;  
store R3 → addr1;
```

Conservatively taint the output if either inputs is tainted

```
??? R4 = load addr2;
```

Can depend on runtime states. Difficult for static analysis to handle

Implicit Information Flow

```
(H) if (secret) {  
(H)   R1 = 10;  
} else {  
(H)   R1 = 20;  
}
```

```
(H) R1 = secret;  
store 10 → R1;
```

Can taint the entire address space!



Taint explosion

Complete Information Flow Tracking from the Gates Up*

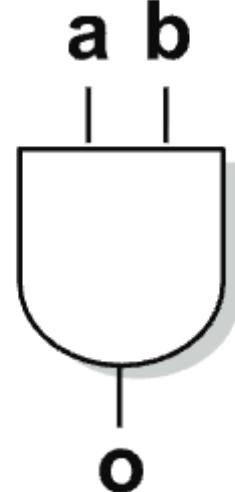
Complete Information Flow Tracking from the Gates Up

Mohit Tiwari Hassan M G Wassel Bita Mazloom Shashidhar Mysore Frederic T Chong
Timothy Sherwood

Department of Computer Science, University of California, Santa Barbara
`{tiwari,hwassel,betamaz,shashimc,chong,she`rwood}@cs.ucsb.edu

*Tiwari et al., "Complete Information Flow Tracking from the Gates Up," ASPLOS '09

A Sound and Precise Taint Propagation Policy for AND Gates



Logic Truth Table

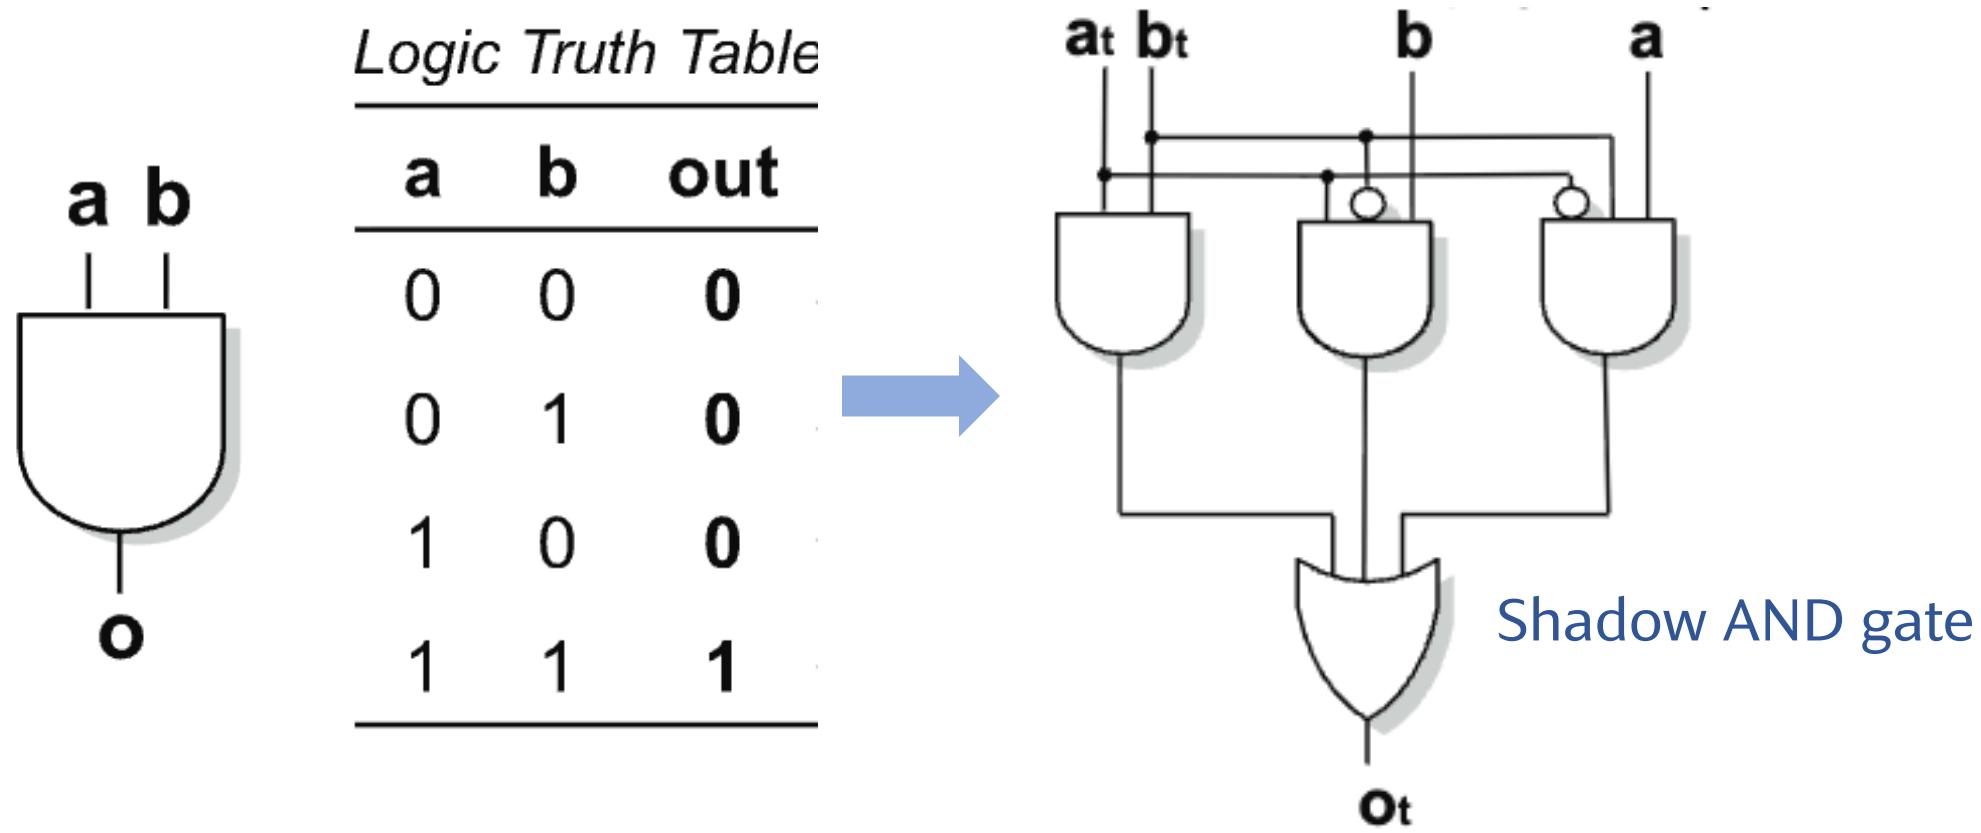
a	b	out
0	0	0
0	1	0
1	0	0
1	1	1

Tainted b and untainted a

a	b	a_t	b_t	out_t
0	0	0	1	0
0	1	0	1	0
1	0	0	1	1
1	1	0	1	1

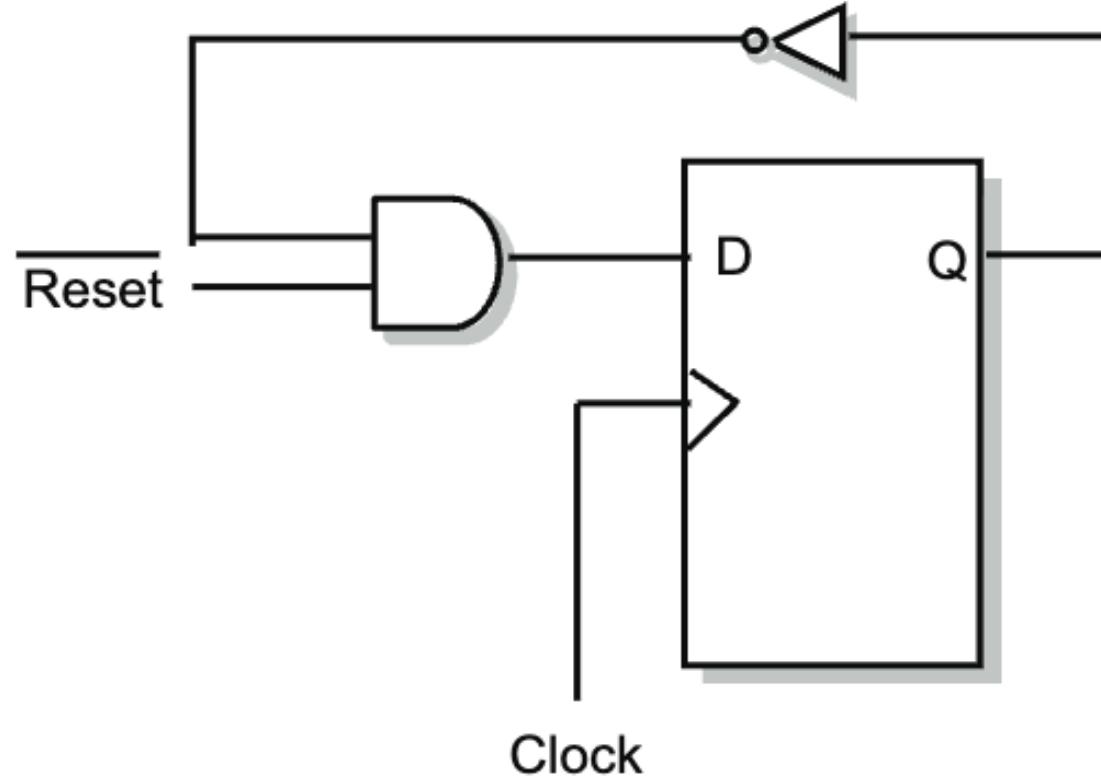
*Tiwari et al., "Complete Information Flow Tracking from the Gates Up," ASPLOS '09

A Sound and Precise Taint Propagation Policy for AND Gates



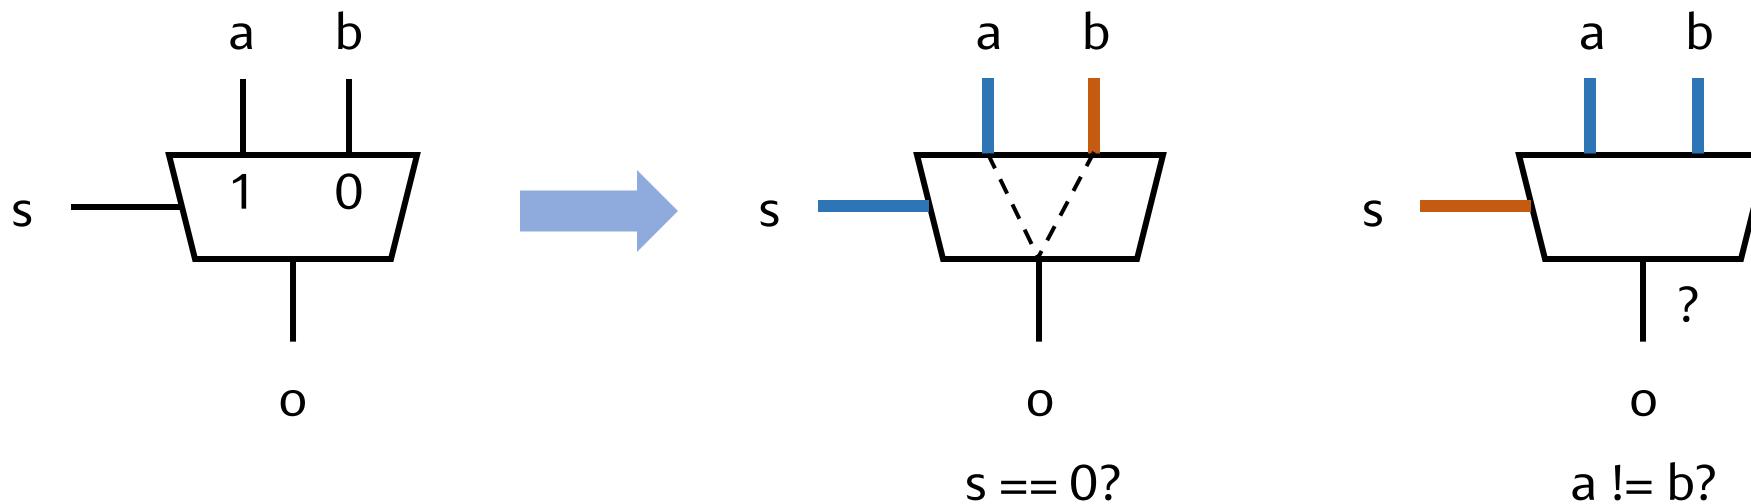
*Tiwari et al., "Complete Information Flow Tracking from the Gates Up," ASPLOS '09

A One-Bit Counter



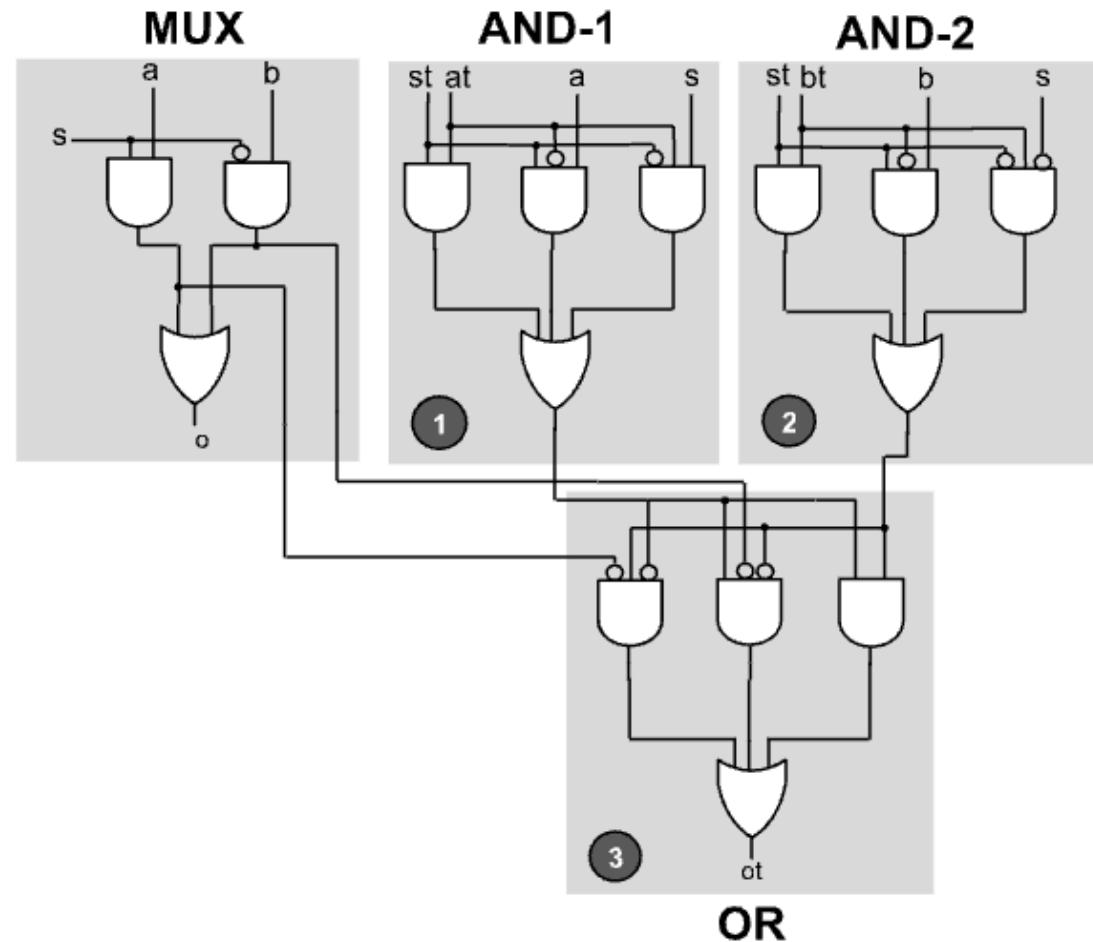
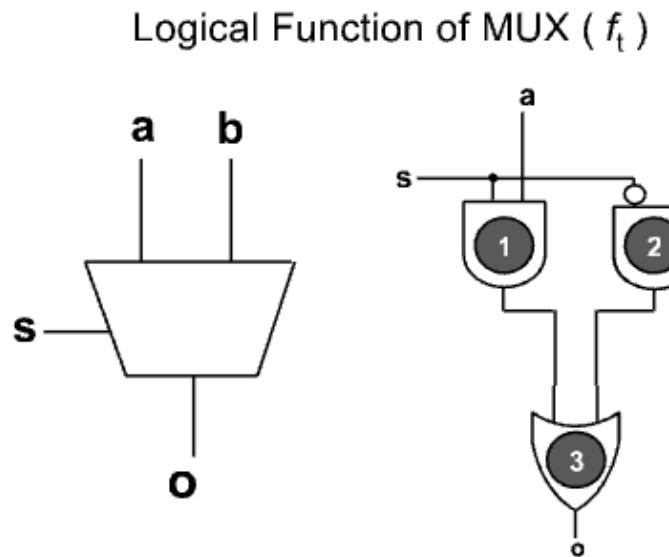
*Tiwari et al., "Complete Information Flow Tracking from the Gates Up," ASPLOS '09

Multiplexer



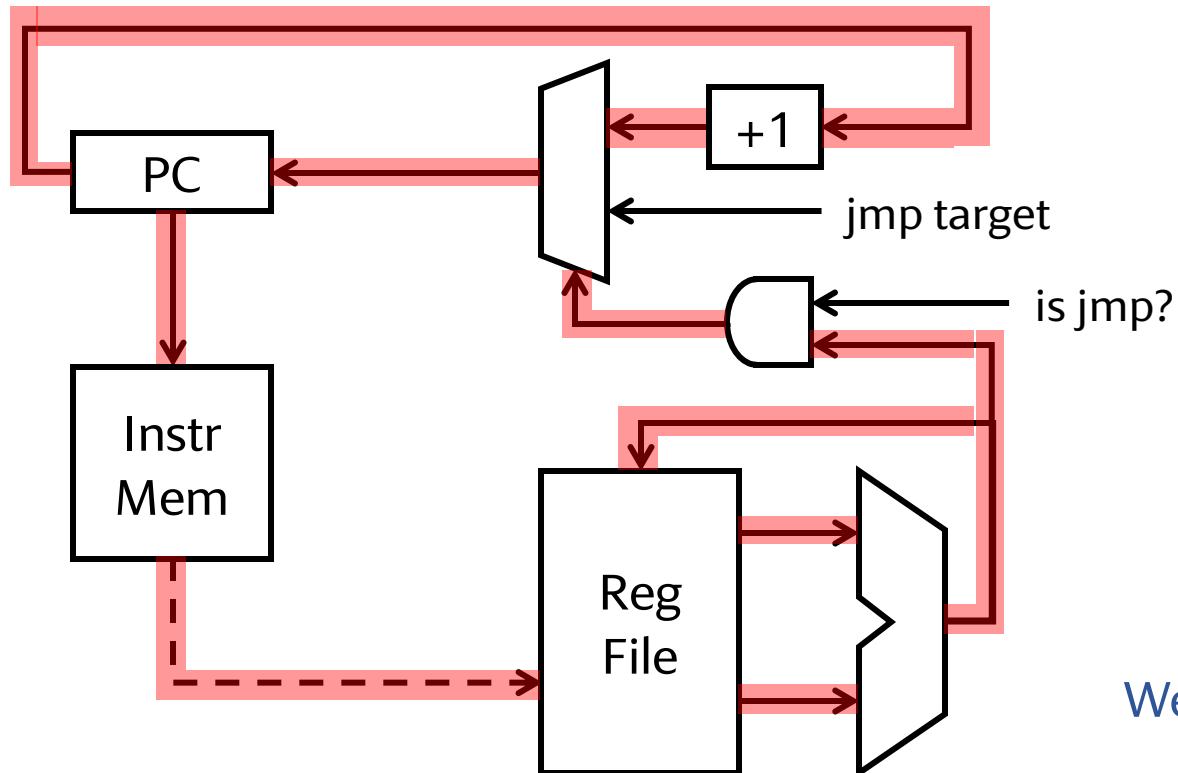
*Tiwari et al., “Complete Information Flow Tracking from the Gates Up,” ASPLOS ’09

Multiplexer



*Tiwari et al., "Complete Information Flow Tracking from the Gates Up," ASPLOS '09

Prevent Taint Explosion: Covert Implicit Flow into Explicit Flow

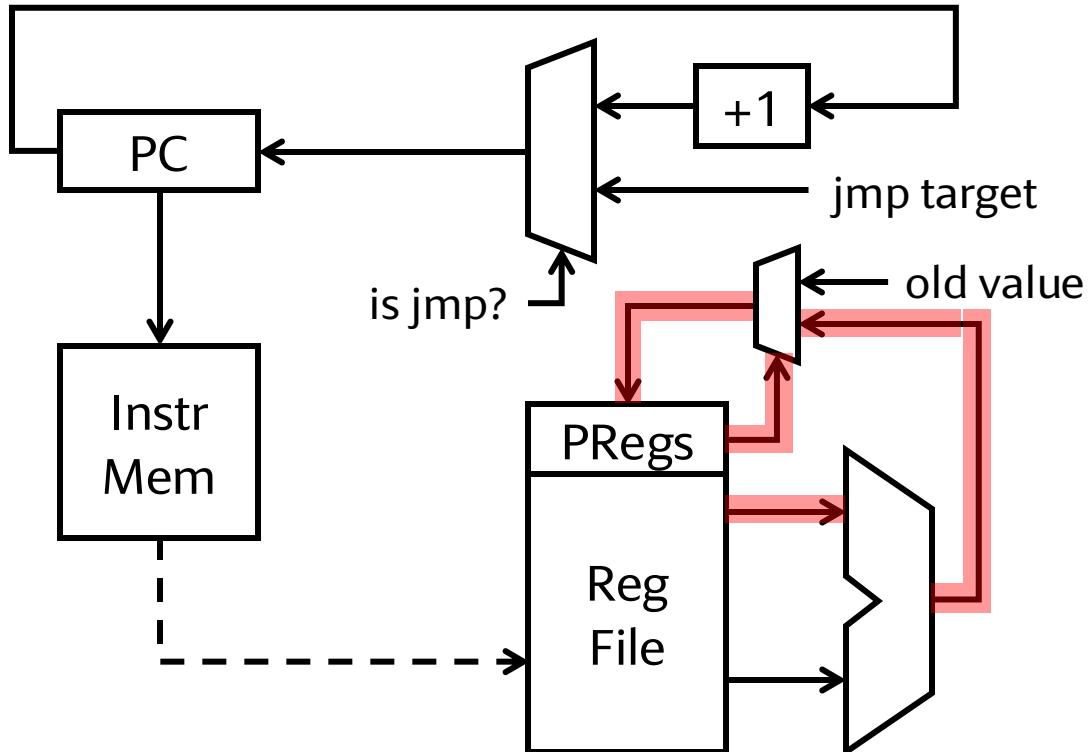


```
0x1: if X == 0; jmp 0x3  
0x2: R1 = 10;  
0x3: R2 = 10;
```

We want to avoid tainting the PC

^{*}Tiwari et al., "Complete Information Flow Tracking from the Gates Up," ASPLOS '09

Prevent Taint Explosion: Covert Implicit Flow into Explicit Flow



0x1: P1 = (X == 0);
0x2: (P1) R1 = 10;
0x3: R2 = 10;

*Tiwari et al., "Complete Information Flow Tracking from the Gates Up," ASPLOS '09

Prevent Taint Explosion: Covert Implicit Flow into Explicit Flow

Bounding Loops

```
0x1: ...
...
0xa: countjump (0x1), 100
```

The number of loop iterations is statically determined

“countjump” cannot be predicated

Constraining Loads/Stores

```
0x1: init-counter C0
0x2: R1 = load 0x200 + C0
0x3: inc-counter C0
0x4: inc-counter C0
...
0xa: countjump (0x2), 100
```

Because the PC cannot be tainted, the addresses are guaranteed to be untainted

*Tiwari et al., “Complete Information Flow Tracking from the Gates Up,” ASPLOS ’09

Putting All Pieces Together

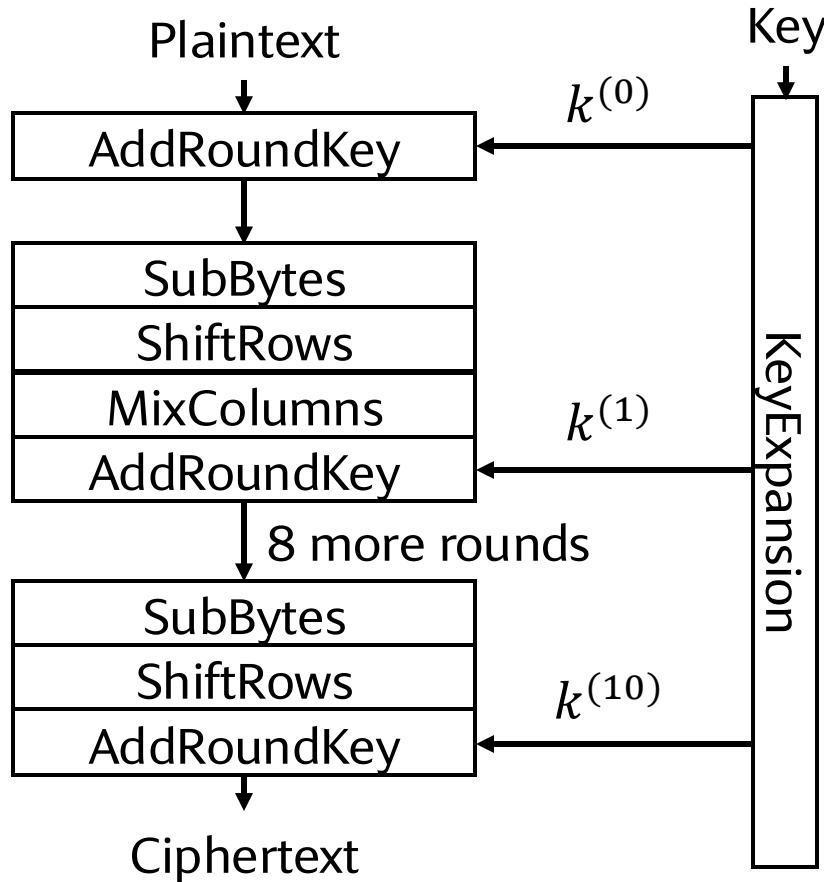
Summary:

- ISA-level restriction that prevents programmer from writing code that can potentially lead to taint explosion
- An FPGA prototype (information-flow tracking at the bit granularity)
- Good for security-sensitive programs
- Area overhead (+70%)
- Performance overhead
- Limited programmability

*Tiwari et al., "Complete Information Flow Tracking from the Gates Up," ASPLOS '09

Speculative Privacy Tracking (SPT)

AES Flow



```
for (size_t r = 1; r <= 10; r++) {  
    eax, ..., edx = AES_Round(eax, ..., edx, rnd_key);  
}  
  
if /* false */ {  
    transmit(eax); // transmitter  
}  
  
Not protected by STT!
```

*Choudhary et al., "Speculative Privacy Tracking (SPT): Leaking Information From Speculative Execution Without Compromising Privacy," MICRO '21

Key Idea of SPT: Leaked in Non-Spec Exec → Not a Secret

```
1 v1, v2 = ...;  
2 transmit(v1);  
3 if /* false */ {  
4     transmit(v1);      ✓  
5     transmit(v1 + 10); ✓  
6     transmit(v1 + v2); ✗  
7 }
```

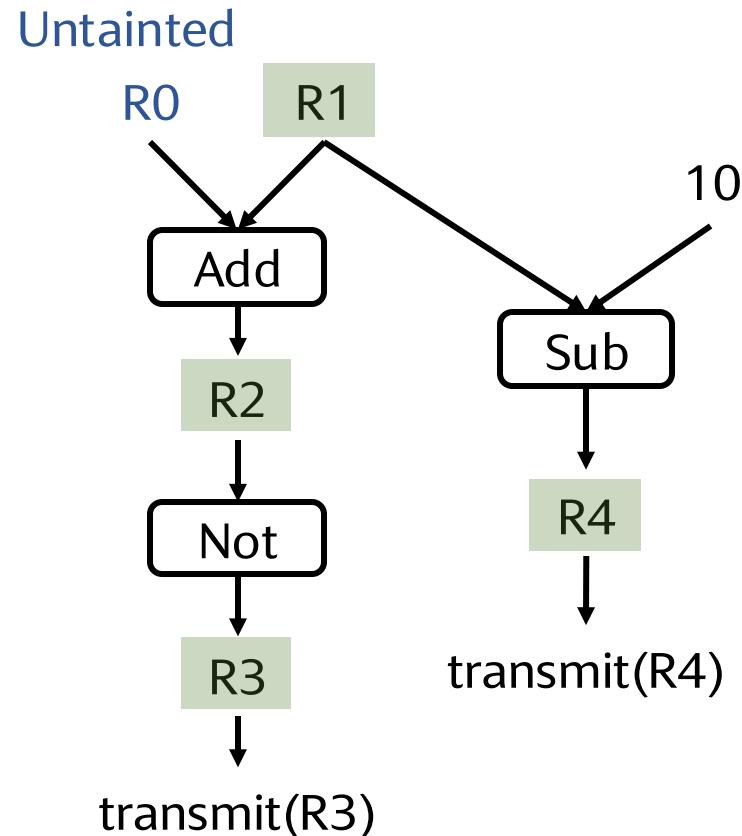
Everything is a secret until it's leaked/declassified by
a non-speculative transmitter

*Choudhary et al., "Speculative Privacy Tracking (SPT): Leaking Information From Speculative Execution Without Compromising Privacy," MICRO '21

Backward Untaint Propagation

```
1 R0, R1 = ...;  
2 R2 = R0 + R1;  
3 R3 = ~R2;  
4 transmit(R3);  
5 if /* false */ {  
6   R4 = R1 - 10;  
7   transmit(R4); ✓  
8 }
```

This example assumes that the program and PC are public



*Choudhary et al., "Speculative Privacy Tracking (SPT): Leaking Information From Speculative Execution Without Compromising Privacy," MICRO '21

Blocking Implicit Channels

Similar to STT, if a branch's predicate is tainted:

- Delay branch resolution
- Delay predictor updates

⇒ PC is not tainted

*Choudhary et al., "Speculative Privacy Tracking (SPT): Leaking Information From Speculative Execution Without Compromising Privacy," MICRO '21

Untaint Propagation Through Memory

```
store R1 → addr1;  
store R2 → addr2;  
store R3 → addr3;  
  
R4 = load addr;
```

```
if (alias(addr, addr3, ...))  
    R4 = R3;  
else if (alias(addr, addr2, ...))  
    R4 = R2;  
else if (alias(addr, addr1, ...))  
    R4 = R1;  
else  
    R4 = load addr;
```

Whether “load addr” issues leaks addr1, addr2, and addr3

*Choudhary et al., “Speculative Privacy Tracking (SPT): Leaking Information From Speculative Execution Without Compromising Privacy,” MICRO ’21

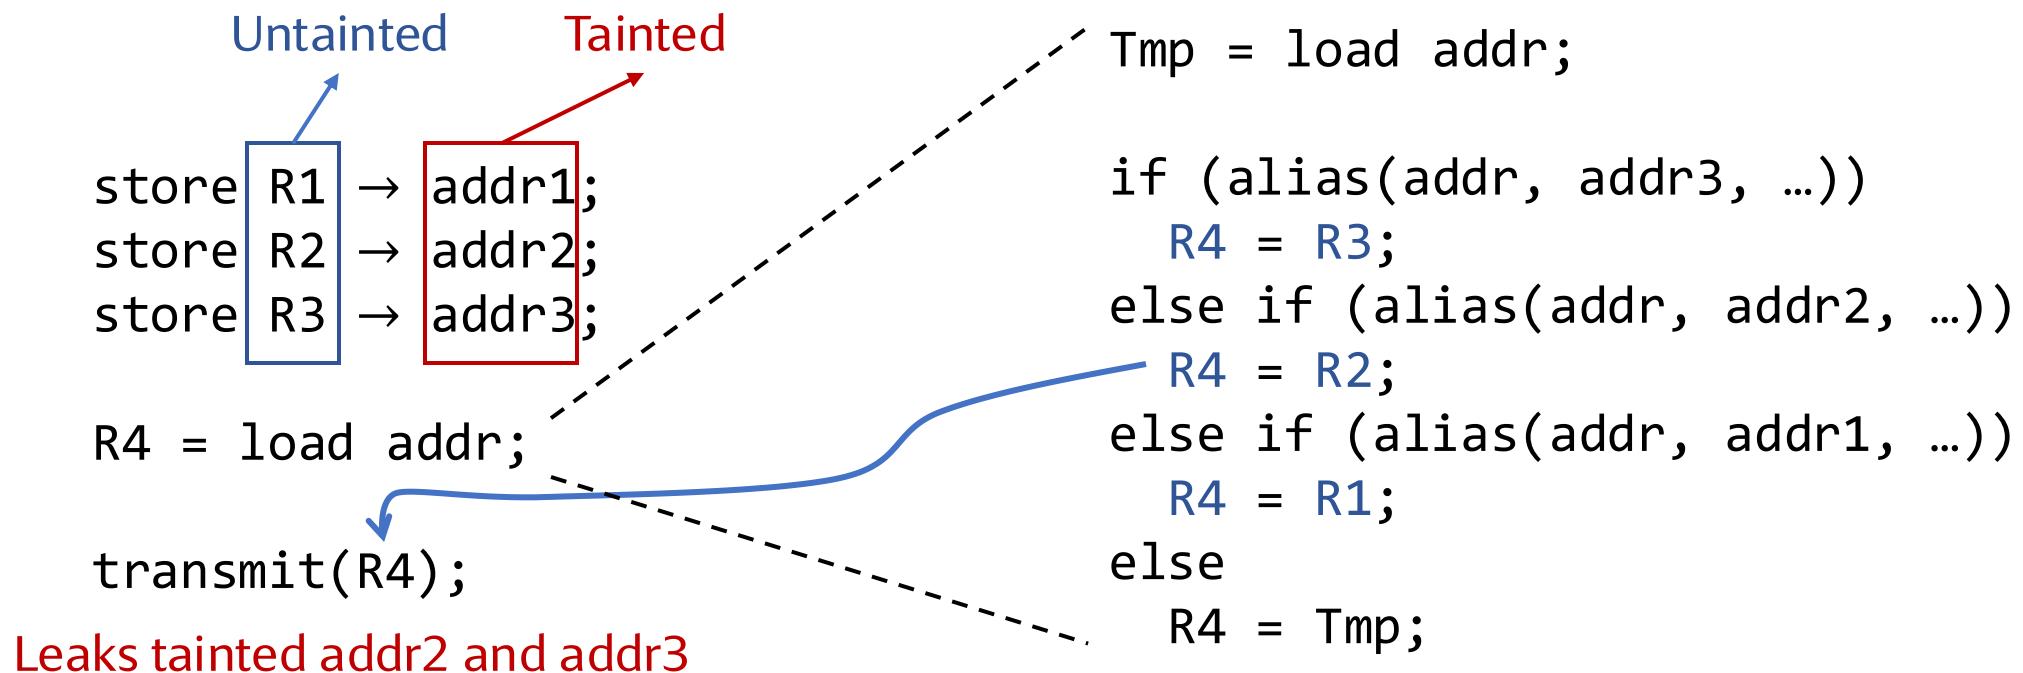
Untaint Propagation Through Memory

```
store R1 → addr1;  
store R2 → addr2;  
store R3 → addr3;  
  
R4 = load addr;
```

```
Tmp = load addr;  
  
if (alias(addr, addr3, ...))  
    R4 = R3;  
else if (alias(addr, addr2, ...))  
    R4 = R2;  
else if (alias(addr, addr1, ...))  
    R4 = R1;  
else  
    R4 = Tmp;
```

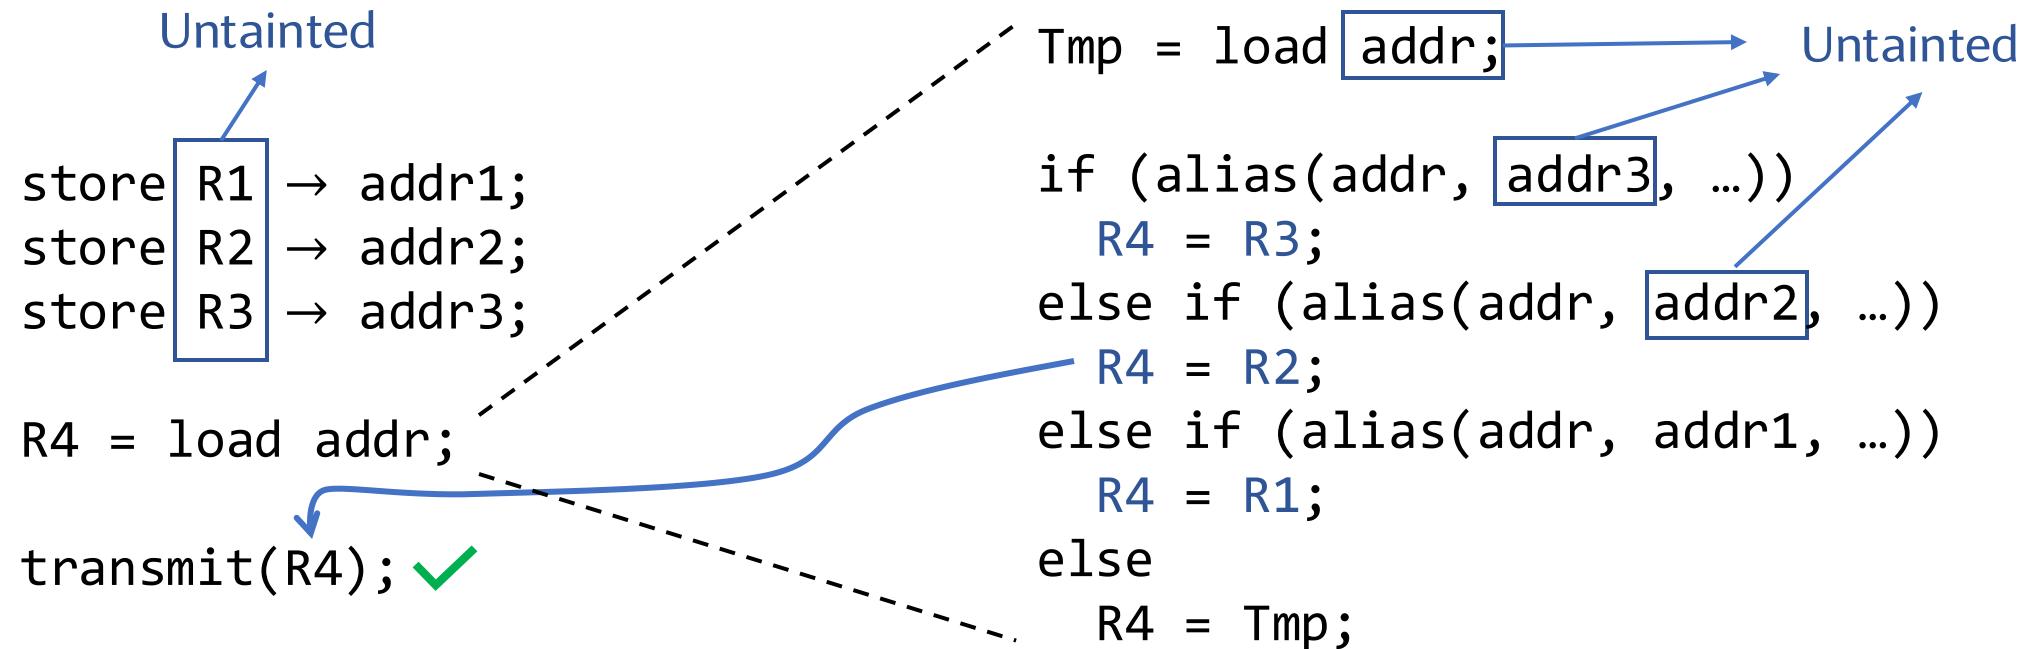
*Choudhary et al., "Speculative Privacy Tracking (SPT): Leaking Information From Speculative Execution Without Compromising Privacy," MICRO '21

Untaint Propagation Through Memory



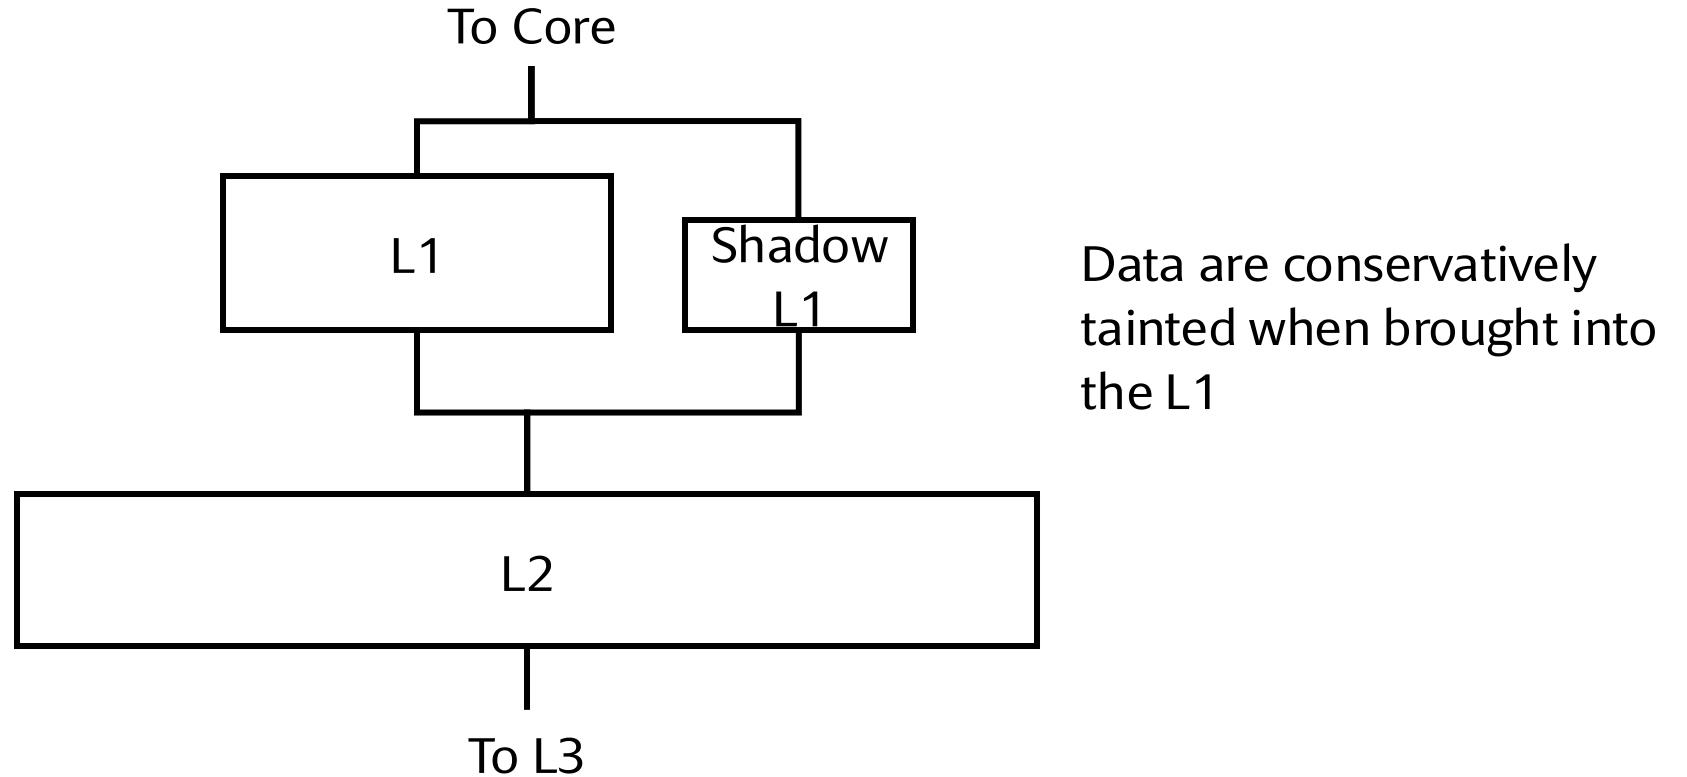
*Choudhary et al., "Speculative Privacy Tracking (SPT): Leaking Information From Speculative Execution Without Compromising Privacy," MICRO '21

Solution: Delay Untaint Propagation in STL Forwarding



*Choudhary et al., "Speculative Privacy Tracking (SPT): Leaking Information From Speculative Execution Without Compromising Privacy," MICRO '21

Untaint Propagation Through Memory (Shadow L1D)



*Choudhary et al., "Speculative Privacy Tracking (SPT): Leaking Information From Speculative Execution Without Compromising Privacy," MICRO '21

Performance Overhead

- gem5 evaluation
- Futuristic model (consider all forms of speculation)
 - $3.6\times$ overhead $\rightarrow 45\%$
- Spectre model (consider only control-flow speculation)
 - $3\times$ overhead $\rightarrow 11\%$
 - STT: 8%

*Choudhary et al., "Speculative Privacy Tracking (SPT): Leaking Information From Speculative Execution Without Compromising Privacy," MICRO '21

Some Questions for Discussion

- [GLIFT] How does GLIFT differ from other tagging/tainting methods discussed earlier in the course, like STT?
- [GLIFT] Could GLIFT serve as the basis for “verifiable secure enclaves” where secrets are provably isolated at the hardware level?
- [GLIFT] Microarchitectural optimization?
- [SPT] Often times paper's like these employ a hardware only or software only solution, if the authors opened up the ability to co-design with software, what kind of solution would SPT now look like? (+ a similar question)
- [SPT] Is there ever any reason to use STT over SPT?
- [SPT] The paper shows 45% overhead in the Futuristic model but only 11% in the Spectre model. Why such a large difference? Which model is more realistic for defending against real-world attacks?